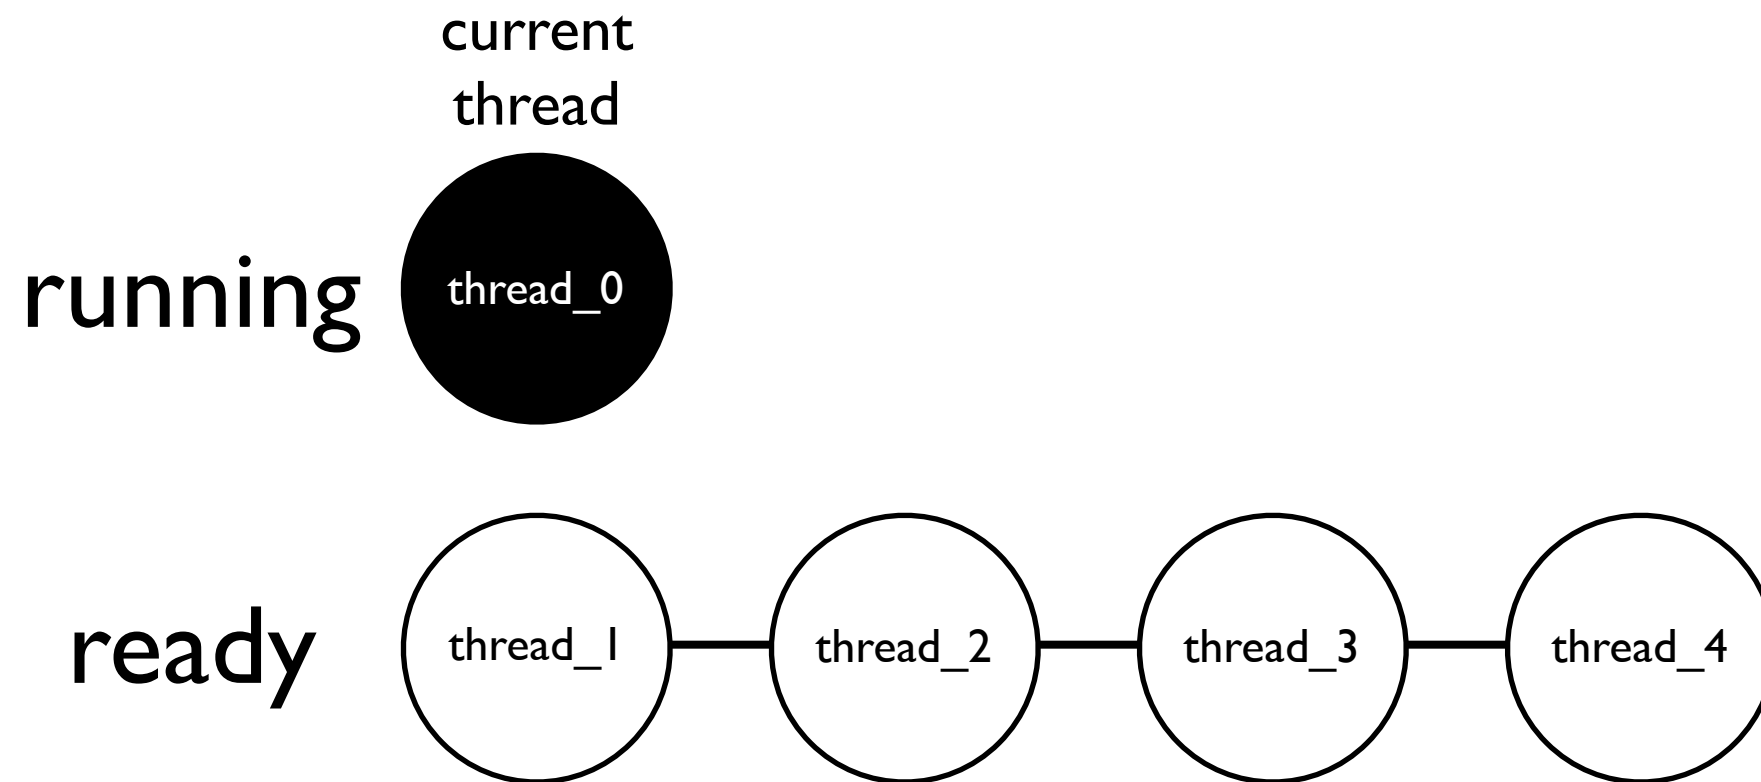


Priority scheduling & Advanced scheduling

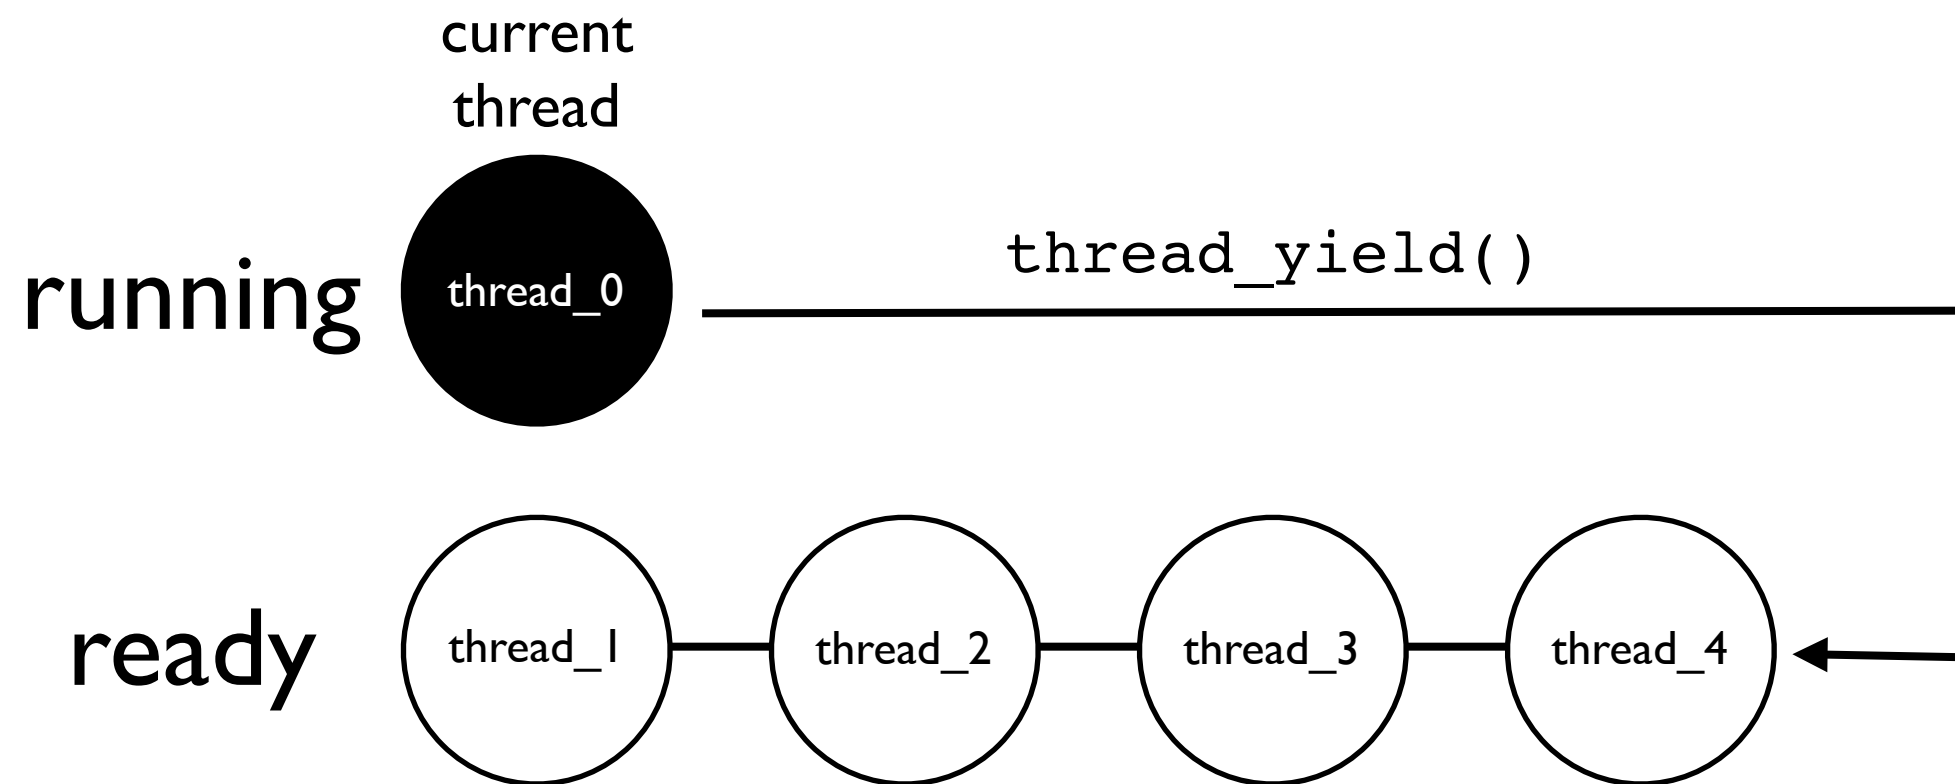
Recap

- So far, we've learnt how to:
 - run and debug pintos tests
 - add our own tests
 - use `lib/kernel/list.h`
 - implement `sleep()` with no busy-wait


Round-robin schedule



Round-robin schedule




Round-robin schedule



```
static struct thread *  
next_thread_to_run (void)  
{  
    if (list_empty (&ready_list))  
        return idle_thread;  
    else  
        return list_entry (list_pop_front (&ready_list), struct thread, elem);  
}
```

Round-robin schedule



```
static struct thread *  
next_thread_to_run (void)  
{  
    if (list_empty (&ready_list))  
        return idle_thread;  
    else  
        return list_entry (list_pop_front (&ready_list), struct thread, elem);  
}
```

Priority scheduling

Priority scheduling

- Some threads/processes require faster response
- This may be achieved with a higher priority
- The next thread to run is that with the highest priority
- If multiple threads have the same highest priority, they alternate in round-robin

Priority scheduling

- A “priority” is an int ranging from 0 to 63
(`PRI_MIN = 0; PRI_DEFAULT = 31; PRI_MAX = 63`)
- A simple priority scheduler in Pintos:
 - Make `next_thread_to_run()` return the thread with highest priority
 - implement `thread_set_priority()`
yield if the new priority is not the highest!
 - make the current thread yield when a higher priority thread is created

Priority scheduling

do you see any
problem with this
scheduler?

Priority scheduling

- Priorities are fixed, unless the user calls `thread_set_priority(...)`
- A cpu-hungry high priority thread may:
 - make the system less responsive (e.g., to I/O)
 - hog the cpu: lower priority threads may starve!

Priority scheduling

- Priorities are fixed, unless the user calls `thread_set_priority(...)`
- A cpu-hungry high priority thread may:
 - make the system less responsive (e.g., to I/O)
 - hog the cpu: lower priority threads may starve!

Solution: advanced scheduler with automatic priorities

Advanced scheduling

(or multilevel feedback queue scheduling - mlfqs)

Advanced scheduling

- The user sets a “nice” value
(ranging from -20 to 20, with 0 as default)
- Priority is automatically adjusted by the OS
(based on the thread’s niceness)
- cpu-hungry threads lose priority
- low-computation threads (like I/O)
automatically get higher priorities
 - the system becomes more responsive to I/O

Advanced scheduling

- To set priorities automatically:
 - each thread has a `nice` value
(again, from -20 to 20)
 - each thread has a `recent_cpu` value
(indicating how much cpu the thread has used recently)
 - global variable `load_avg` stores the system load
(in number of ready+running threads per second in the last minute)
- To adjust each thread's priority:
$$\text{priority} = \text{PRI_MAX} - (\text{recent_cpu} / 4) - (\text{nice} * 2)$$

Advanced scheduling

nice (per thread): type int
initial value: set by user
user can change with `thread_set_nice()`

load_avg (global): type fixed-point real
initial value: 0
every second:
$$\text{load_avg} = (59/60) * \text{load_avg} + (1/60) * \text{ready_or_running_threads}$$

recent_cpu (per thread): type fixed-point real
initial value: 0
every tick: increase `recent_cpu` by 1 for current thread
every second:
$$\text{recent_cpu} = (2 * \text{load_avg}) / (2 * \text{load_avg} + 1) * \text{recent_cpu} + \text{nice}$$

priority (per thread): type int
initial value: calculated based on the initial value of `nice`
every 4 ticks:
$$\text{priority} = \text{PRI_MAX} - (\text{recent_cpu} / 4) - (\text{nice} * 2)$$

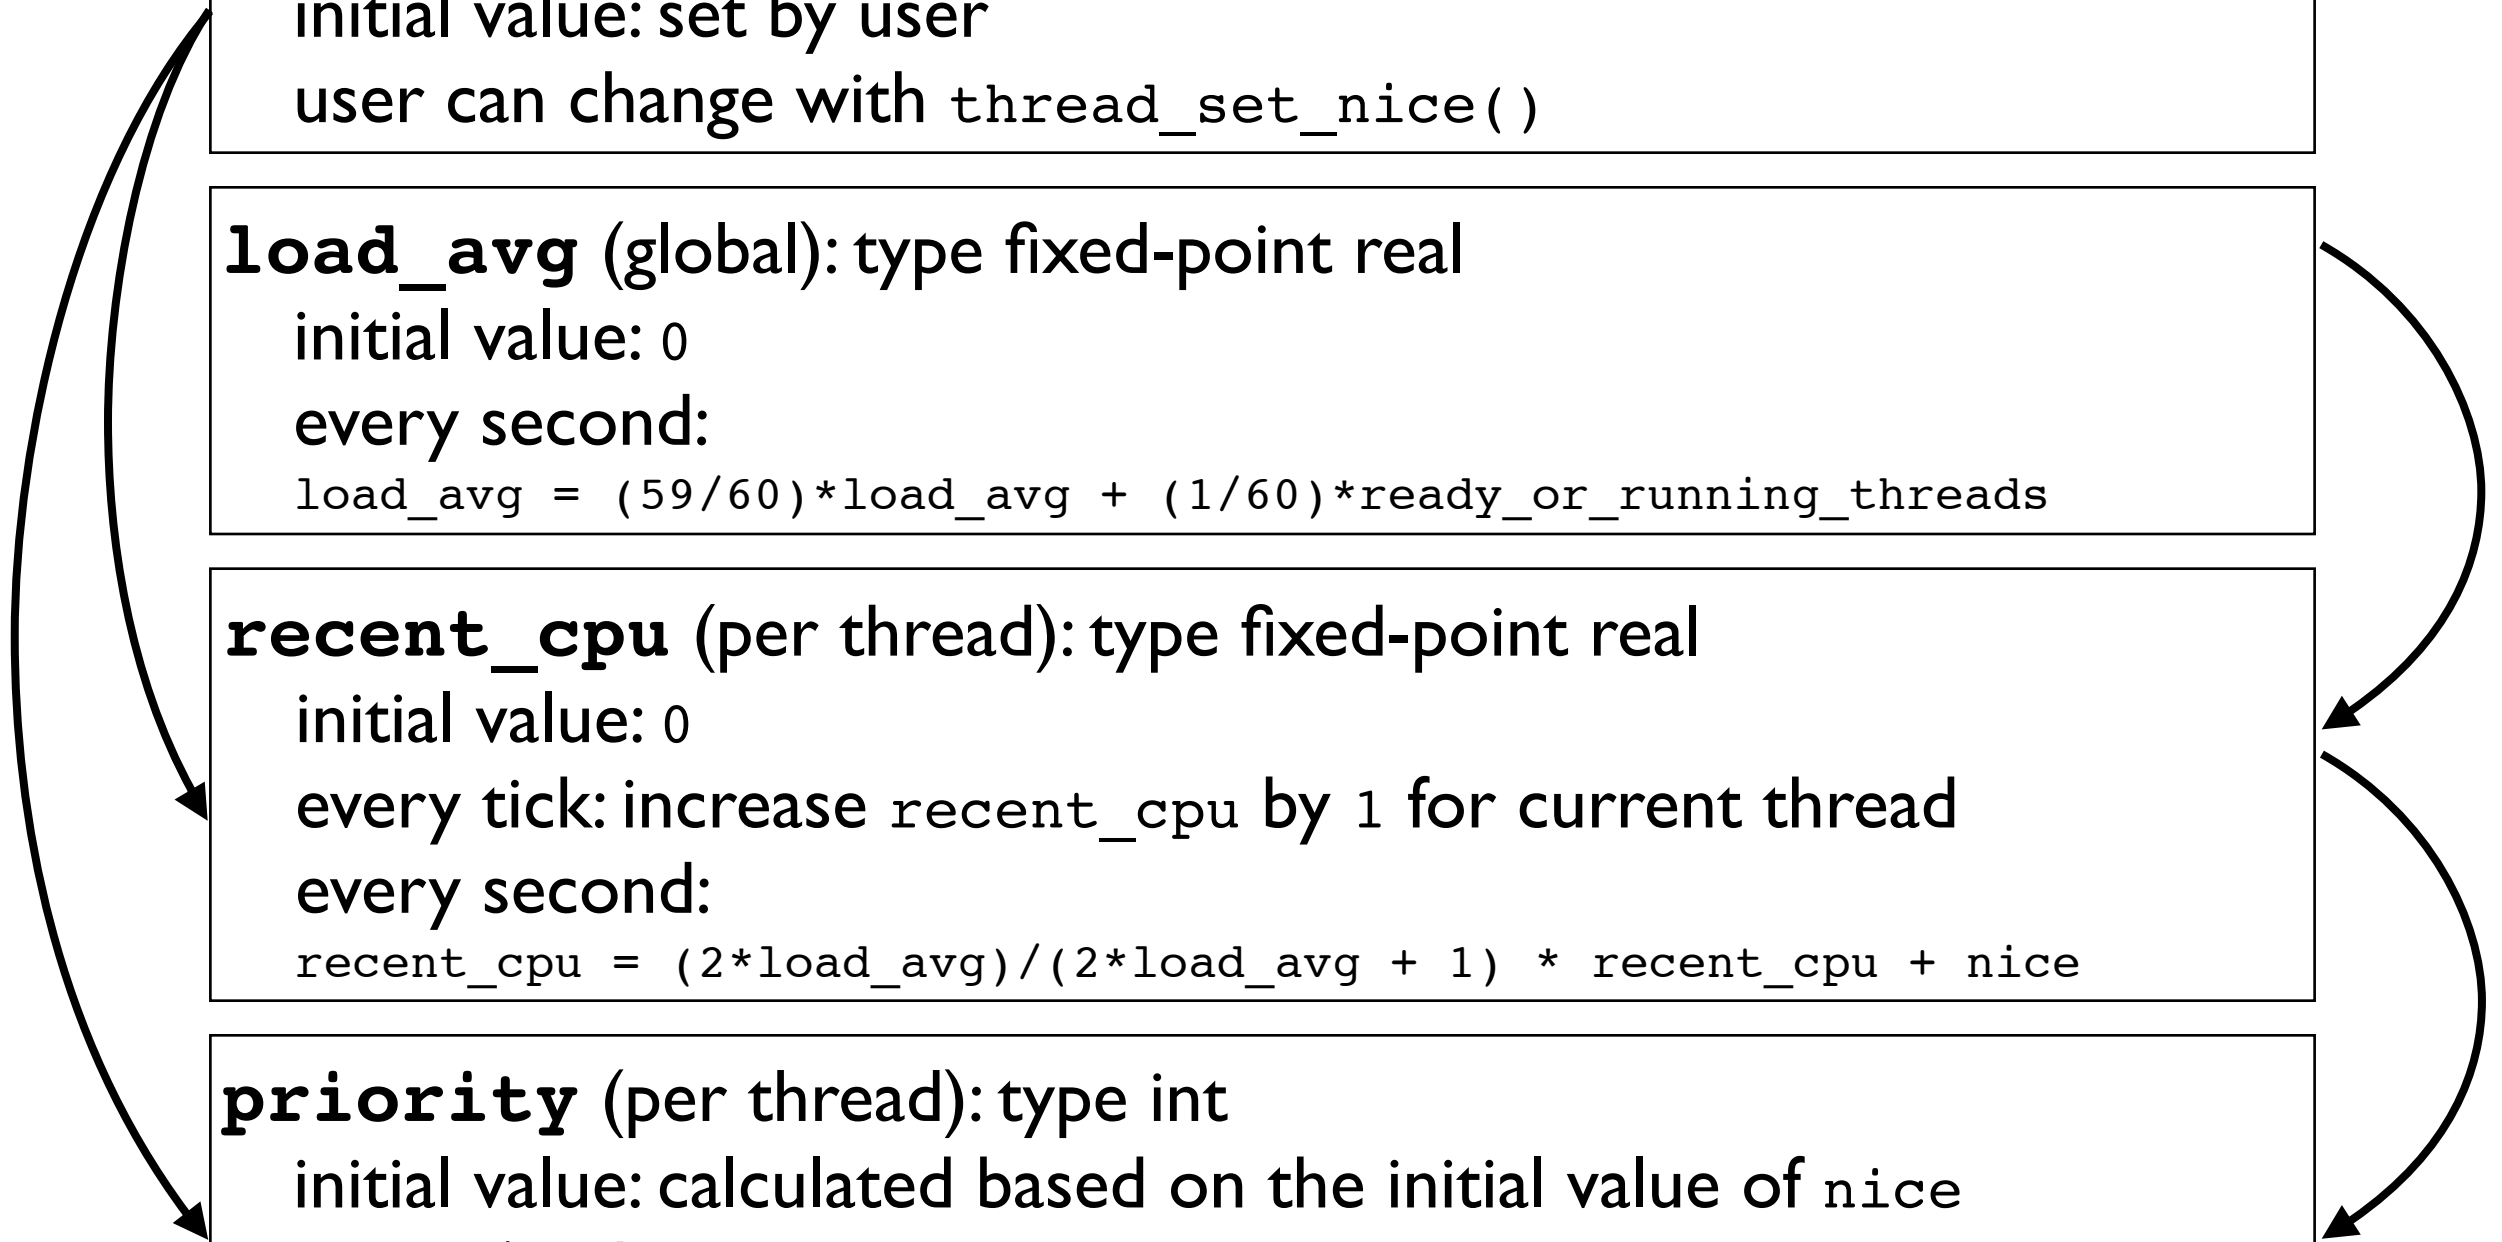
Advanced scheduling

nice (per thread): type int
initial value: set by user
user can change with `thread_set_nice()`

load_avg (global): type fixed-point real
initial value: 0
every second:
$$\text{load_avg} = (59/60) * \text{load_avg} + (1/60) * \text{ready_or_running_threads}$$

recent_cpu (per thread): type fixed-point real
initial value: 0
every tick: increase `recent_cpu` by 1 for current thread
every second:
$$\text{recent_cpu} = (2 * \text{load_avg}) / (2 * \text{load_avg} + 1) * \text{recent_cpu} + \text{nice}$$

priority (per thread): type int
initial value: calculated based on the initial value of `nice`
every 4 ticks:
$$\text{priority} = \text{PRI_MAX} - (\text{recent_cpu} / 4) - (\text{nice} * 2)$$



Fixed-Point Real

- `load_avg` and `recent_cpu` are real numbers
- Usually, no float operations in the kernel
- Fixed-point real variables are a solution
 - Implemented using integers
 - You can use the header `fpr_arith.h`, available on moodle

Advanced scheduling

- Functions that must be implemented:
 - `thread_set_nice(int)`:
update the thread nice, recalculate its priority based on that and yield if the new priority is not the highest anymore
 - `thread_get_nice()` : return the thread's nice value
 - `thread_get_load_avg()` : return `load_avg * 100`
 - `thread_get_recent_cpu()` :
return current thread's `recent_cpu * 100`
- The current thread must yield if a thread with higher priority (based on its `nice` value) is created

Advanced scheduling

- Check the bool variable `thread_mlfqs`
 - It tells whether the option `-mlfqs` was passed to the kernel
 - If true, use the advanced scheduler
 - If false, use the simple priority scheduler

Assignment

- In class: implement the simple priority scheduler
- Home: extend your scheduler with the advanced scheduling
 - it must be possible to select which one, by giving (or omitting) the -mlfqs option

Assignment

```
$ make check
```

```
...
```

```
pass tests/threads/alarm-single
```

```
pass tests/threads/alarm-multiple
```

```
pass tests/threads/alarm-simultaneous
```

```
pass tests/threads/alarm-priority
```

```
pass tests/threads/alarm-zero
```

```
pass tests/threads/alarm-negative
```

```
pass tests/threads/priority-change
```

```
FAIL tests/threads/priority-donate-one
```

```
FAIL tests/threads/priority-donate-multiple
```

```
FAIL tests/threads/priority-donate-multiple2
```

```
FAIL tests/threads/priority-donate-nest
```

```
FAIL tests/threads/priority-donate-sema
```

```
FAIL tests/threads/priority-donate-lower
```

```
pass tests/threads/priority-fifo
```

```
pass tests/threads/priority-preempt
```

```
FAIL tests/threads/priority-sema
```

```
FAIL tests/threads/priority-condvar
```

```
FAIL tests/threads/priority-donate-chain
```

```
pass tests/threads/mlfqs-load-1
```

```
pass tests/threads/mlfqs-load-60
```

```
pass tests/threads/mlfqs-load-avg
```

```
pass tests/threads/mlfqs-recent-1
```

```
pass tests/threads/mlfqs-fair-2
```

```
pass tests/threads/mlfqs-fair-20
```

```
pass tests/threads/mlfqs-nice-2
```

```
pass tests/threads/mlfqs-nice-10
```

```
FAIL tests/threads/mlfqs-block
```

```
10 of 27 tests failed.
```

Readings

- Read the Pintos documentation
- Chapter 2
 - 2.2.3
 - 2.2.4
 - 2.3
- Appendix B